

MULTI-AGENT WORKFLOW EXPERIMENT

Routing AI agents in 0.2 seconds for a fraction of a cent

12 / 12

clear requests
routed right

~30×

faster than the LLM
step

≈ \$14

per million
decisions

Dispatcher

Jev · one choice

1 OF 3

IT claims

Human
resources

100% sure · 0.16 s

Devices
management

Workflow built with the **AI Assistant from SynergyAI** · Decisions by **Jev** (TypeSafe AI)

Letting an AI that can't write text route our agents: 0.2 seconds and a fraction of a cent per decision

In multi-agent workflows, one step quietly costs time and money every time it runs: **deciding which agent should handle the request**. We tried a new kind of AI model for that step. It never writes a sentence. It only picks from the options you give it. Here is what we found, including the part that surprised us.

12 / 12

clear requests sent to the
right team, across two
workflows

~0.16 s

median decision time, vs.
~5 s for the LLM step it
replaces

≈ \$14


per million routing
decisions (342 input
tokens each, output free)

0

chance of naming a team
that isn't on the map:
answers come only from
your list

The workflow under test

The workflow below was built with the **AI Assistant from SynergyAI**. It is a typical internal help desk. A request comes in, a "dispatcher" agent called techBuddy reads it, and sends it to **exactly one** of three teams: IT claims, Human resources or Devices management. The "1 OF 3" badge on the canvas shows that rule: three branches are drawn and only one runs. HR requests continue into an HR playbook connected to Google Calendar, Slack, Workday, HiBob and UKG.



The workflow, designed with SynergyAI's AI Assistant. Every agent here runs on Claude. The experiment replaced only one thing: techBuddy's routing decision.

Think of the dispatcher as a receptionist. Normally a large language model (LLM) does this job. It reads the request, thinks it over in text, and names a team. That works, but it takes about 5 seconds of an 11-second run. Sometimes it also names a team that doesn't exist, so the software needs safety nets to catch those answers.

What is Jev?

Jev is a "System One" model from **TypeSafe AI**, released in limited early access on 15 September 2026. It answers closed questions: which of these options, what level on this scale, or is this true. It returns a probability for each option, and nothing else. It cannot write text, so it cannot invent an answer that is not on your list.

That makes it a natural fit for a dispatcher. The question "which of these three teams owns this request?" is a multiple-choice question.

Test 1: three clearly different teams

We sent seven requests to the real help desk dispatcher.

Request	Jev chose	Confidence
Two days of vacation next week	Human resources ✓	100%
Account locked out	IT claims ✓	100%
Cracked laptop screen	Devices management ✓	100%
VPN drops every ten minutes	IT claims ✓	100%
Vacation days left this year	Human resources ✓	100%
Need a second monitor	Devices management ✓	100%
Something is wrong, I need help	IT claims ✗	100%

Six out of six were right, at a median of 159 milliseconds. The last request has no clue in it at all, yet Jev sent it to IT with full confidence. A simple "hello" went to IT at 97%.

This is not a bug. It is what the model promises to do. Offered three teams and nothing else, it must pick one of the three. A rule like "only trust answers above X% confidence" would never have caught it.

The fix: make "I can't tell" one of the options

If the model can only answer with options you supply, then uncertainty has to be one of those options. We added a fourth choice, **"unclear"**, to every question.

Request	Three teams only	With "unclear" added
Something is wrong, I need help	IT 99% ✗	unclear 100% ✓
hello	IT 97% ✗	unclear 100% ✓
I need access	IT 100% ✗	unclear 69% / IT 31%
Laptop can't reach the VPN	IT 97% ✓	IT 98% ✓

Request	Three teams only	With "unclear" added
Leaving, returning my laptop	Devices mgmt 99% ✓	Devices mgmt 98% ✓
Two days of vacation	HR 100% ✓	HR 100% ✓

Real requests still route with confidence. Empty ones now say so. "I need access" is genuinely ambiguous, and it came back split: 69% unclear, 31% IT. The probabilities only became informative once uncertainty had somewhere to go.

Test 2: two teams that overlap

A second workflow chooses between **Billing** and **Sales**, which overlap in real life.

Request	Jev chose	Confidence	Split
Charged twice this month	Billing	100%	Billing 100%
Refund on my subscription	Billing	100%	Billing 100%
Invoice shows the wrong plan	Billing	97%	Billing 98%
How much is the Pro plan	Sales	99%	Sales 100%
A demo before we commit	Sales	100%	Sales 100%
Discount for non-profits	Sales	97%	Sales 98%
Upgrade to the annual plan	Sales	49%	Sales 66% / Billing 33%
Monthly vs annual, and when charged	Billing	58%	Billing 72% / Sales 16%
A question about my account	unclear	99%	unclear 100%
hello	unclear	100%	unclear 100%


Every clear request scored 95% or more. Questions like "Upgrade to the annual plan" scored below 60%, split between Sales and Billing. That is the right reading: Sales sells the upgrade, Billing takes the money. Without the "unclear" option, "hello" would have gone to Sales.

So the two kinds of problem are caught by two different signals:

"unclear" means the request names no team. Hand it to a human.

Confidence below 60% means two teams could own it. Let the LLM weigh it.

The design we shipped: fast when sure, careful when not



Jev decides only when it is sure. The LLM keeps the cases that need judgment. "Unclear" requests go to a person. A guess never decides a route.

Every outcome other than a confident answer falls back to the existing LLM path: no API key, a timeout, an error, an unexpected answer. **The fast path can make routing cheaper. It cannot make it fail.**

What this means if you run AI workflows

Measure	Jev	LLM routing step
Time per decision (median)	159 ms	about 5,000 ms
Time range	134 to 385 ms (30 calls)	not measured
Tokens per decision	342 in, 39 out	a full agent turn
Billing	input only, \$0.042 per million	input and output
Cost per decision	about \$0.0000144	not recorded
Can answer outside the list	no	yes, needs safety nets

Cost

Billed on input tokens only, at \$0.042 per million. One decision used 342 tokens: about **\$0.0000144**, roughly **\$14 per million** decisions. The LLM's token cost for the same step wasn't recorded, so no precise ratio is claimed.

Speed

Median **159 ms** (134–385 ms over 30 calls) vs. **~5 s** for the LLM step. Routing was ~5 s of an 11 s run, so the whole run gets noticeably faster. The other agents' time stays the same.

Consistent decisions

Answers are always one of your options: an off-menu answer is **impossible**. With an "unclear" option and a 60% threshold, clear and ambiguous cases separated in every test. The catch: a number, not a reason, so a wrong route would be silent.

Fits into a workflow

Nothing new to write. Team descriptions are the agents' own instructions. Jev sits in front of the existing LLM step, so nothing is removed and the other agents (Claude-based here) are unchanged.

Availability

Limited **early access**, one vendor, one API key. Weak with numbers, dates and adversarial text, per the vendor. Every failure falls back to the LLM, so an outage only **slows** routing down. It doesn't stop it.

The lesson

Give the model a way to say **"I don't know"**, and keep a careful path for the cases in between.

What we don't know yet

Small sample. 30 calls on two workflows with well-named teams. The 60% threshold hasn't been tested on six similar teams or on messy, real-world wording.

No hand-over notes. When the LLM routes, it also writes a short note for the chosen agent. Jev can't. Nothing breaks, but we haven't measured whether answers get worse. That is the next experiment.

Black box. For routes where a wrong branch is expensive, the evidence isn't strong enough yet.

Vendor claims are the vendor's. TypeSafe advertises 40 to 200 times faster than frontier LLMs. We measured about 30 times faster on this step. Independent benchmarks are still scarce.

The takeaway

For a narrow, closed decision like routing a request, a model that can only choose is fast, very cheap and predictable. The key lesson isn't really about speed: **give the model a way to say "I don't know"**, and keep a more careful path for the cases in between. Next

we will test whether the same idea works for yes/no checks, like approval gates, in agent workflows.

Are you using small, specialised models for decisions inside your agent workflows? I'd like to hear what has worked for you and what hasn't.

#AIAgents #MultiAgentSystems #WorkflowAutomation #LLM #AICosts #SynergyAI #Jev

Sources: experiment report "Jev for dispatcher routing: experiments, design, results" (27 September 2026, model jev-1.13.0). Background on Jev:

<https://www.datacamp.com/blog/system-one-models-jev> and

<https://pasqualepillitteri.it/en/news/16460/typesafe-jev-chatless-ai-beats-claude>